

# Introduction to Deep Learning

## Lecture 7 Generative Models and GANs

Maria Vakalopoulou & Stergios Christodoulidis

**MICS Laboratory**  
**CentraleSupélec**  
**Université Paris-Saclay**



CentraleSupélec


Wednesday, January 10, 2024



# Last Lecture


# Recurrent Neural Networks (RNNs)

- They were first introduced in 1986.
- They are neural networks with loops in order to allow information to persist.
- These loops represent the influence of previous value on the same value at the current step.
- For a simpler representation we could unroll this RNN in time.



# LSTMs

- LSTMs are an RNN architecture.
- They are capable of learning long-term dependencies
- Tackling the vanishing gradient problem
- The LSTM has been found extremely successful in many applications



(<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>)

# Transformers

- A sequence-to-sequence model with an Encoder-Decoder architecture.
- Even though it is not a recurrent network, it is designed such that it can work with sequences
  - It utilizes positional encoding of the input to capture the relative positions across
  - It utilizes a self-attention mechanism to decide which other parts of the sequence are important.
- The state-of-the-art for a lot of applications, extended to vision
  - Vision Transformers

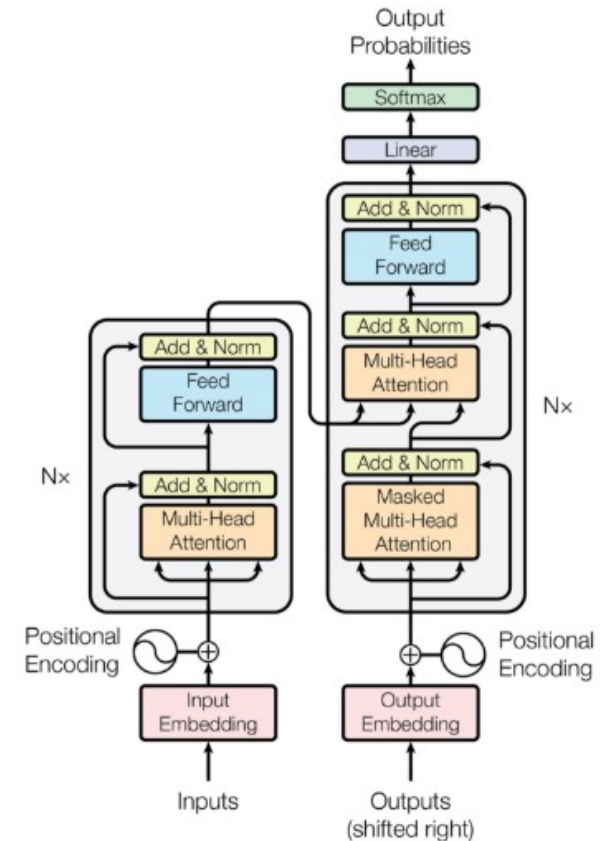
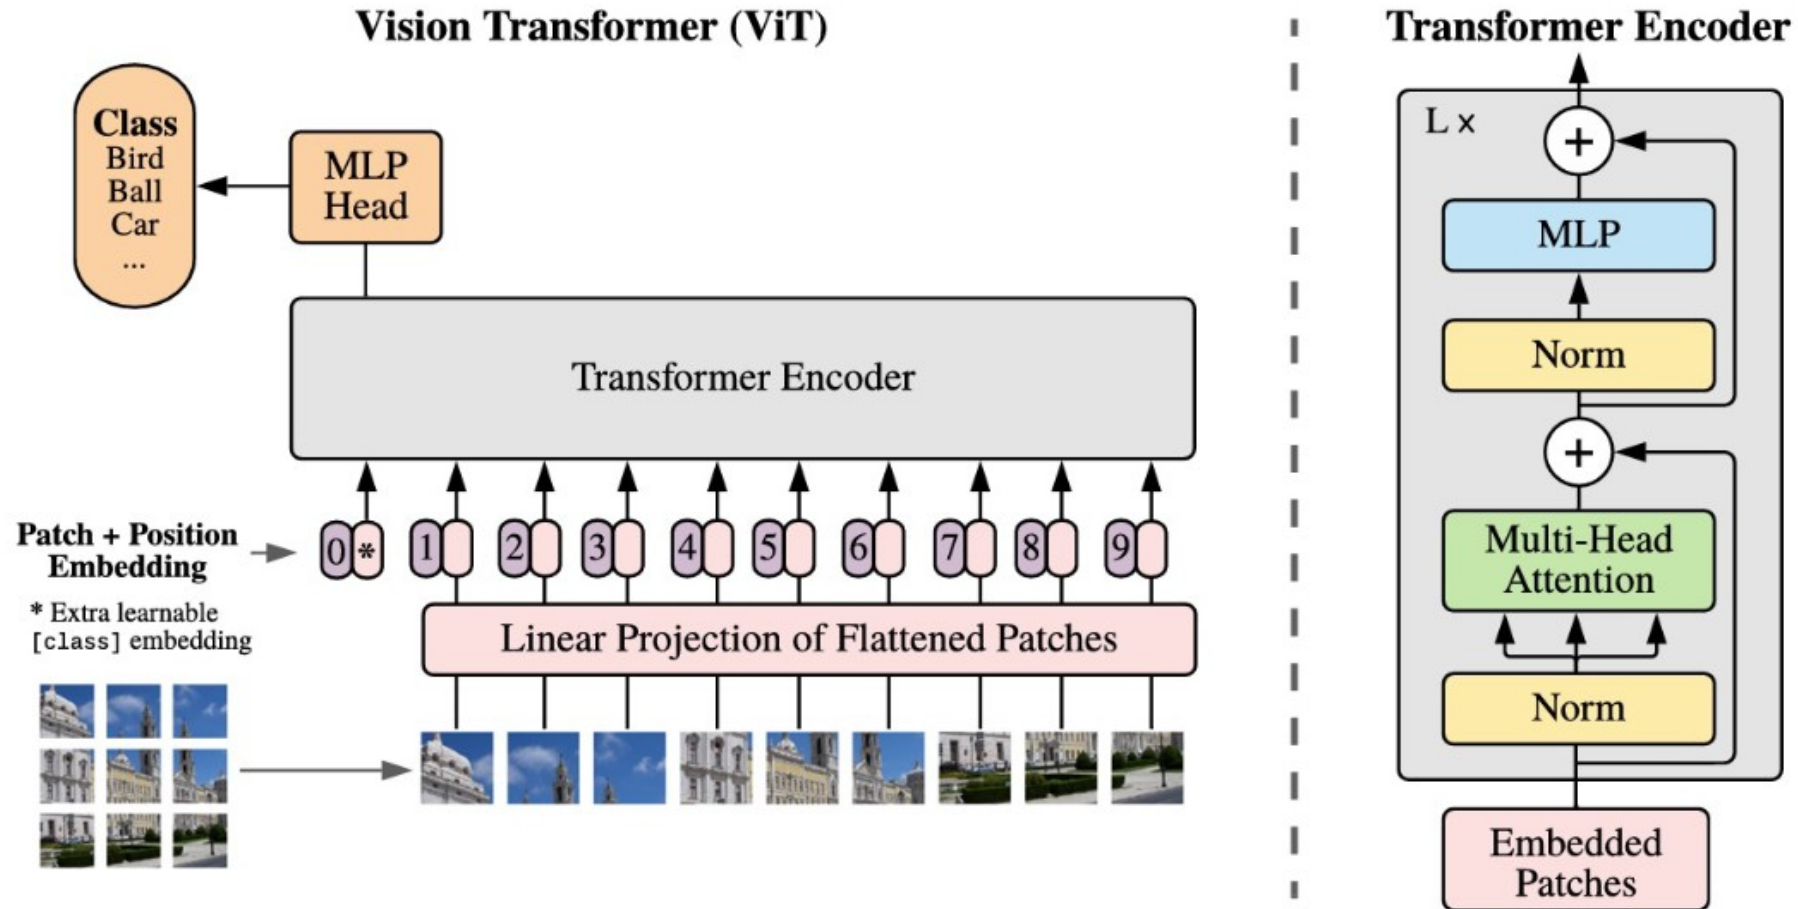


Figure 1: The Transformer - model architecture.

# Vision Transformers



Dosovitskiy, Alexey, et al. "An image is worth 16x16 words: Transformers for image recognition at scale." arXiv preprint arXiv:2010.11929 (2020).

# Today's Lecture

# Today's Lecture

- Gentle intro to generative models
- Variational Autoencoders
- Generative Adversarial Networks
- Variants of Generative Adversarial Networks



# Types of Learning

- Generative modelling
  - Learn the joint pdf:  $p(x,y)$
  - Model the world → Perform tasks, e.g. use Bayes rule to classify:  $p(y|x)$
  - Naive Bayes, Variational Autoencoders, GANs



# Types of Learning

- Generative modelling
  - Learn the joint pdf:  $p(x,y)$
  - Model the world  $\rightarrow$  Perform tasks, e.g. use Bayes rule to classify:  $p(y|x)$
  - Naive Bayes, Variational Autoencoders, GANs
- Discriminative modelling
  - Learn the conditional pdf:
  - Task-oriented
  - E.g., Logistic Regression, SVM



# Types of Learning

- What to pick?
  - V. Vapnik: “One should solve the *[classification]* problem directly and never solve a more general *[and harder]* problem as an intermediate step”
- Typically, discriminative models are selected to do **the** job
- Generative models give us more theoretical guarantees that the model is going to work as intended
  - Better generalization
  - Less overfitting
  - Better modelling of causal relationships

# Applications of generative modeling?

# Applications of generative modeling?

- Act as a regularizer in discriminative learning
  - Discriminative learning often too goal-oriented
  - Overfitting to the observations
- Semi-supervised learning
- Simulating “possible futures” for Reinforcement Learning
- Data-driven generation/sampling/ simulation

# Applications of generative modeling?

- Image Generation



(a) Generated by LSGANs.



(b) Generated by DCGANs (Reported in [13]).



2014



2015



2016



2017



2018



# Applications of generative modeling?

- Super-resolution


original



bicubic  
(21.59dB/0.6423)



SRResNet  
(23.44dB/0.7777)



SRGAN  
(20.34dB/0.6562)



# Applications of generative modeling?

- Cross-model translation

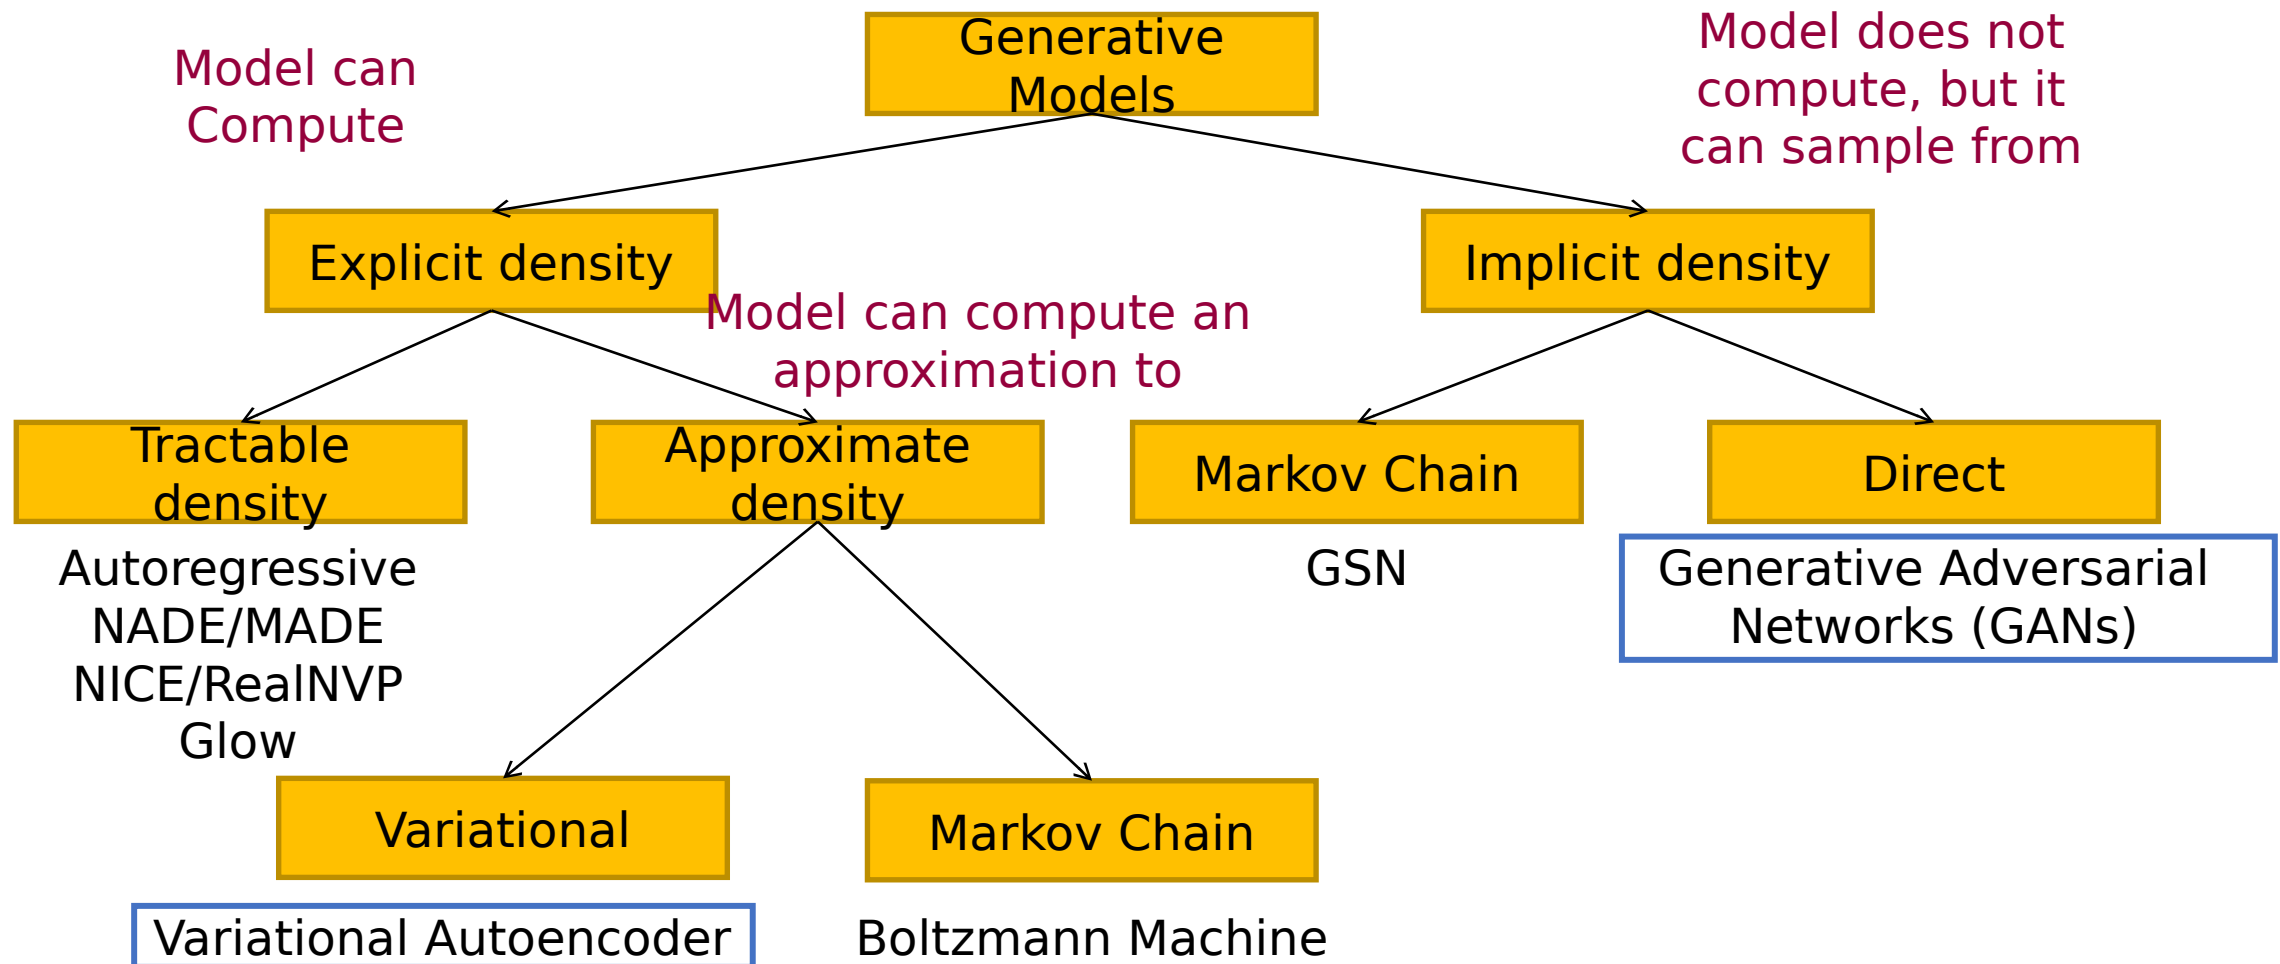





# Other Applications

- M. Mustafa et al. “Cosmogan: Creating High-Fidelity Weak Lensing Convergence Maps Using Generative Adversarial Networks” in Arxiv 2017
- S. Collaboration. “Fast Simulation of Muons Produced at the SHiP Experiment Using Generative Adversarial Networks”. In Arxiv 2019
- Z. E. et al. “Deep Learning Enables Rapid Identification of Potent DDR1 Kinase Inhibitors” In. Nature Biotechnology 2019
- Deep Fakes.

# A map of generative models




# Variational Autoencoders




# Latent Representation Models

- Latent variables are high level features:
  - In combination they generate the data
- In Latent Representation models we are interested to learn these representations
  - i.e., Figure out what these latent variables are.
  - Can we learn these in a supervised manner?




# Autoencoders (AE)

- Idea: Use the latent representation to reconstruct the input data (Autoencoding == encoding itself)
- Encoder part (Green): maps the observed data  $x$  to a latent representation  $z$
- Decoder part (Yellow): reconstructs the observed data using the latent representation  $z$




# Variational Autoencoders (VAE)

- Generative spin on autoencoders
- Can we use the model to generate data?
- There is no explicit way to know how to set  $z$




# Variational Autoencoders (VAE)

- Instead of straight learning the latent representation  $z$
- We learn the parameters of a multivariate gaussian from which we sample  $z$
- Not deterministic any more → Stochastic sampling operation



# Variational Autoencoders (VAE)


- Probabilistic twist:
  - Encoder learns  $p(z|x)$
  - Decoder learns  $p(x|z)$  **Generative Model!**





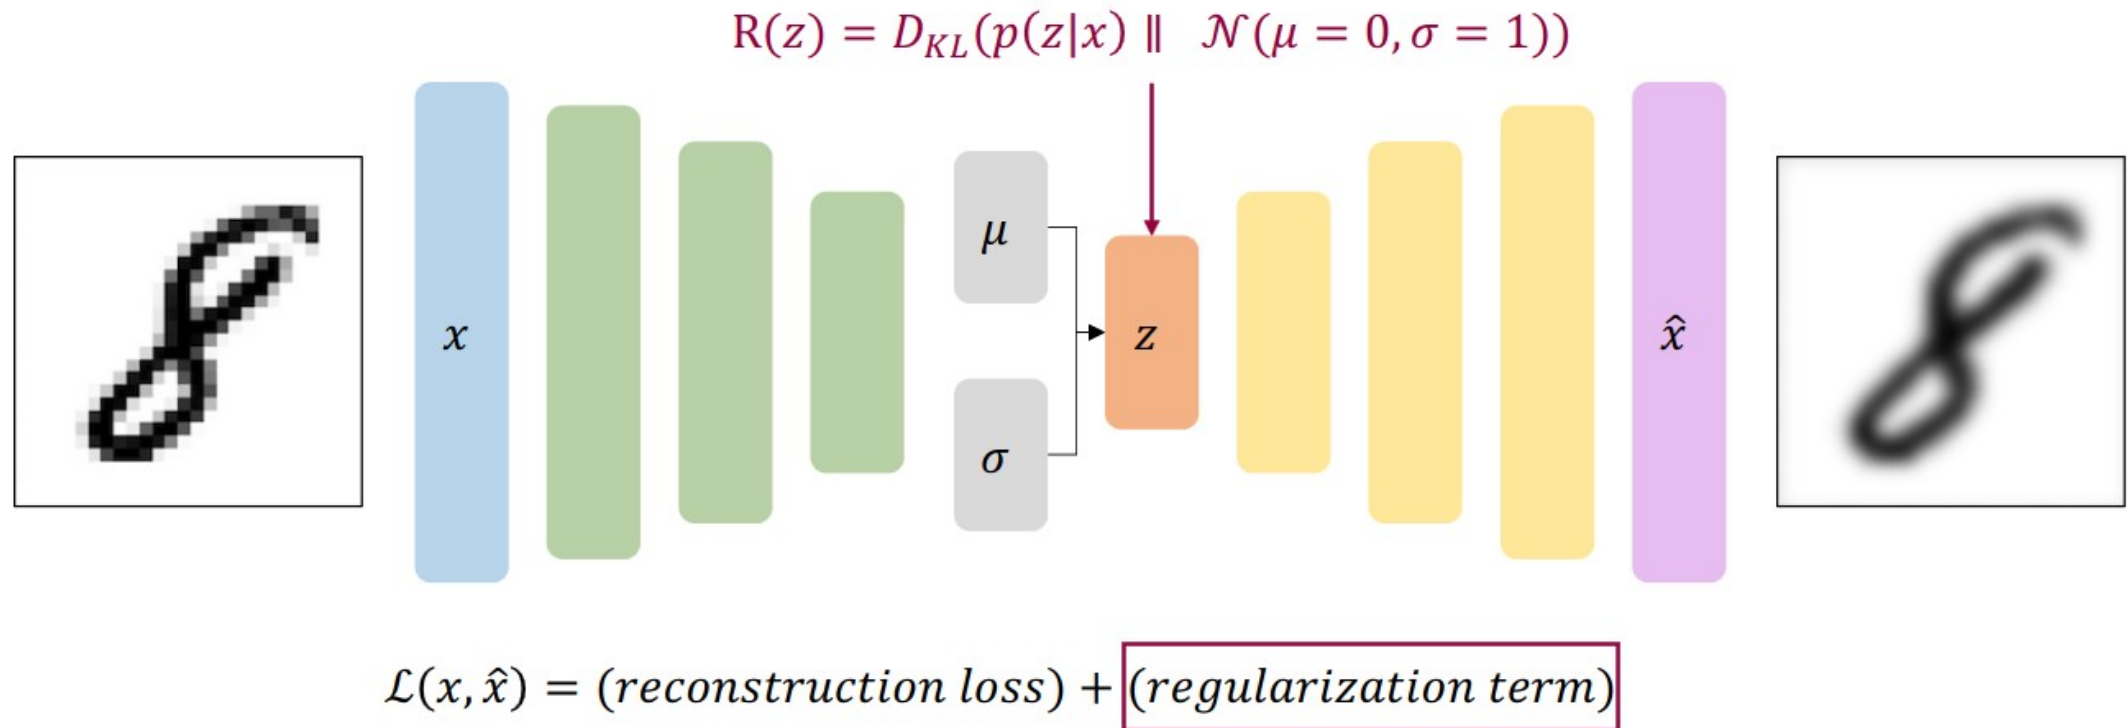
# Variational Autoencoders (VAE)

- Reconstruction loss
  - Mean Square Error




# Variational Autoencoders (VAE)

- Regularization Term
  - KL divergence between the inferred latent distribution and a prior distribution.
  - Typically zero mean unit std normal distribution



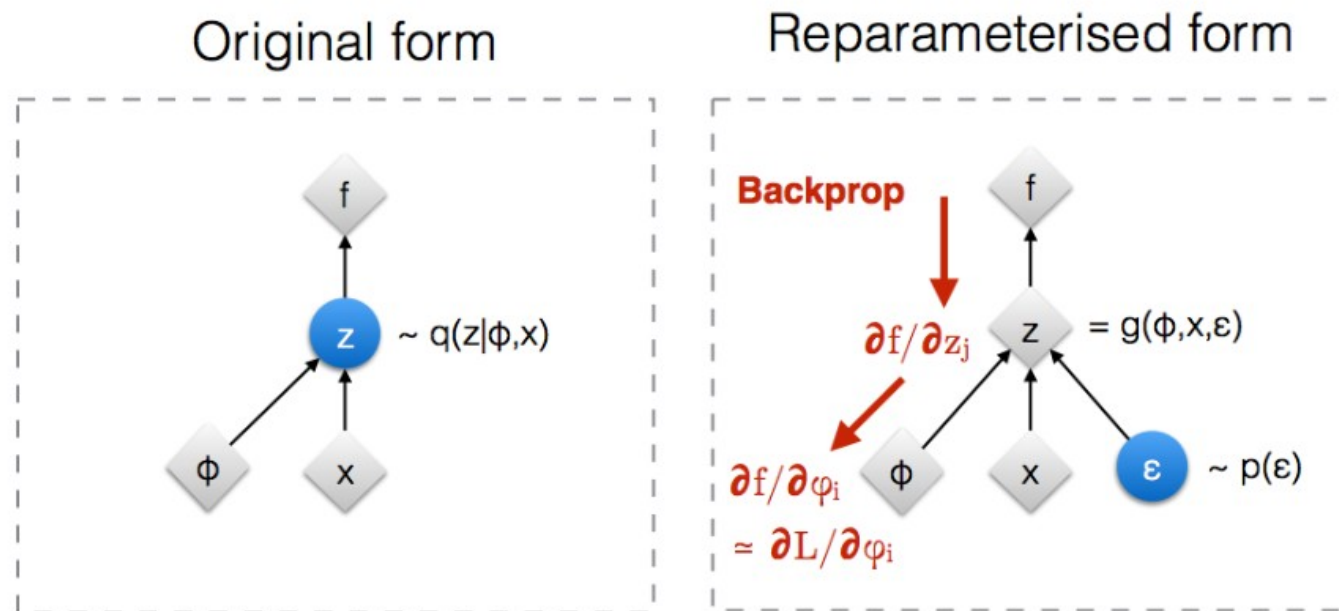
# Variational Autoencoders (VAE)

- Sampling operation is not differentiable
  - How can we solve this?



$$\mathcal{L}(x, \hat{x}) = (\text{reconstruction loss}) + (\text{regularization term})$$

# Reparameterization Trick



◊ : Deterministic node

● : Random node

[Kingma, 2013]  
[Bengio, 2013]  
[Kingma and Welling 2014]  
[Rezende et al 2014]

# VAE: Summary

- Compresses observed data to some smaller representation
- Unsupervised setting
- Reparametrization trick for end-to-end training
- Force latent representation to imitate a gaussian distribution (KL divergence)
- The latent variables can be interpreted by perturbation their values
- Can be used to generate new samples




# GANs



# What is a GAN?

- **G**enerative
  - You can sample novel input samples
  - E.g., you can literally “create” images that never existed
- **A**dversarial
  - Our generative model **G** learns adversarially, by fooling an discriminative oracle model **D**
- **N**etwork
  - Implemented typically as a (deep) neural network
  - Easy to incorporate new modules
  - Easy to learn via backpropagation

# Adversarial Learning




Goodfellow, 2019



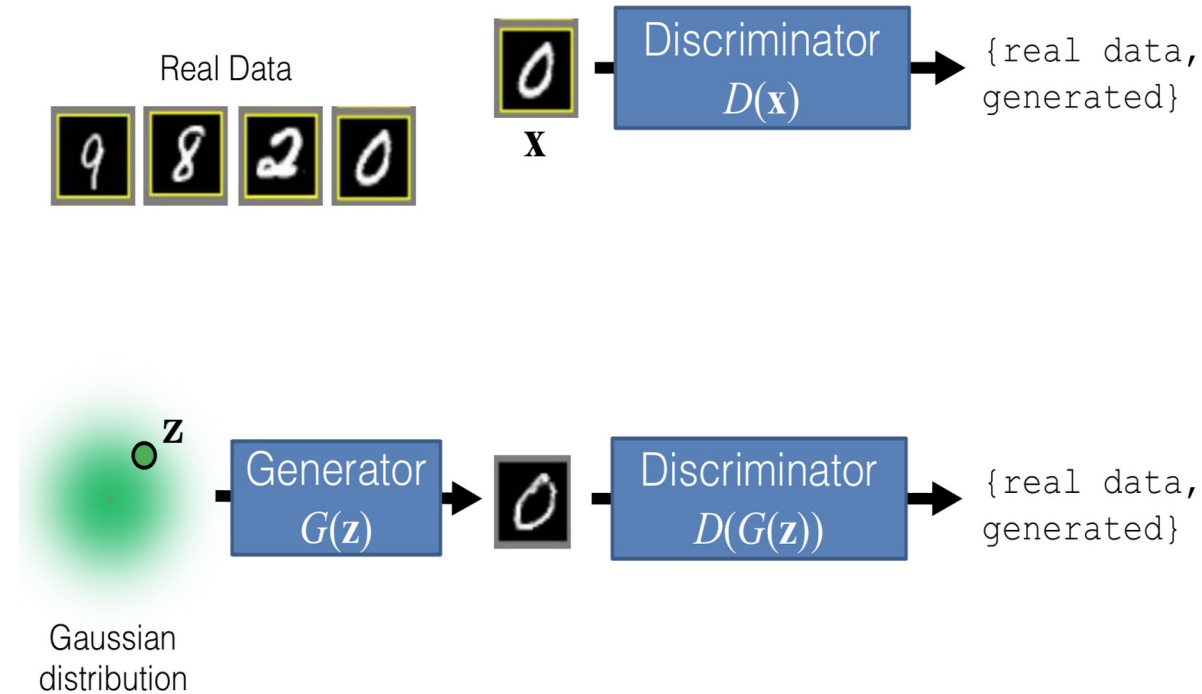
# Generative Adversarial Networks

- We would like to train a network  $G$  to generate images from some domain from random vectors  $z$ :

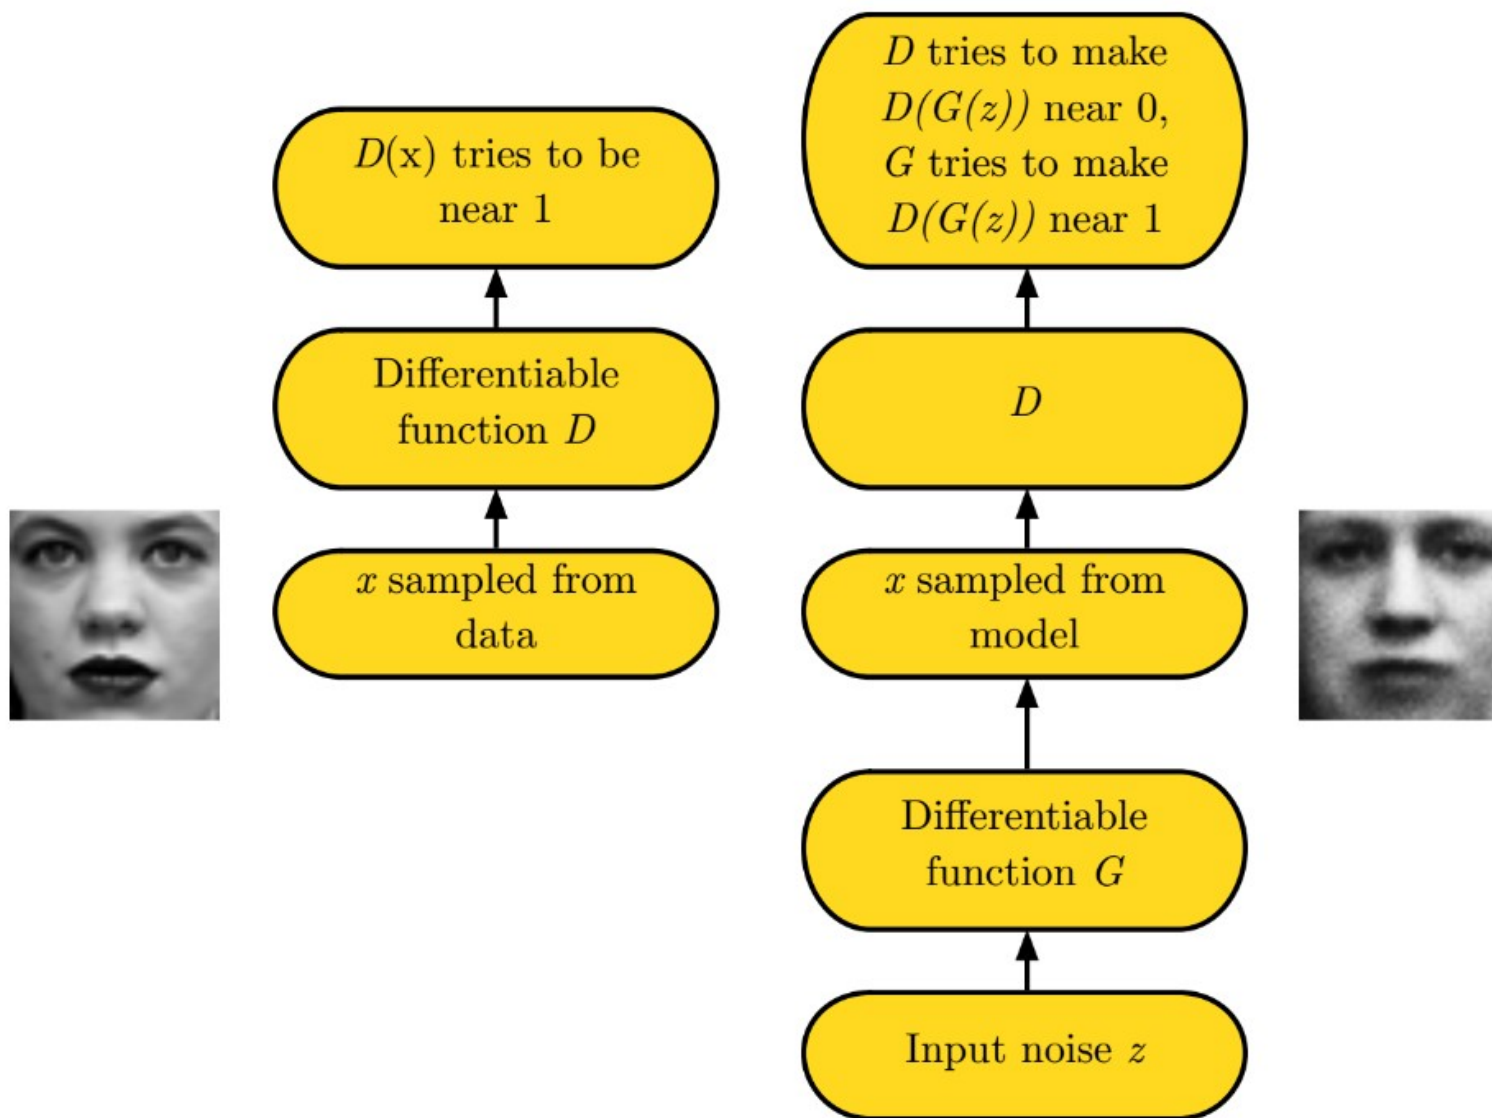


# Generative Adversarial Networks

- Idea: Add a second network (Discriminator  $D$ ) jointly trained with the Generator  $G$  to recognize if an input is a real sample from the domain of interest or if it was created by the Generator.
- When the Discriminator cannot distinguish the generated images from the real ones, the Generator generates realistic images




# GAN: Pipeline



# Generator network $x = G(z; \theta^{(G)})$

- Must be differentiable
- No invertibility requirement
- Trainable for any size of  $z$
- Can make conditionally Gaussian given  $z$ , but no strict requirement



# Generator & Discriminator: Implementation

- The discriminator is just a standard neural network
- The generator looks like an inverse discriminator





Figure 2. Our text-conditional convolutional GAN architecture. Text encoding  $\varphi(t)$  is used by both generator and discriminator. It is projected to a lower-dimensions and depth concatenated with image feature maps for further stages of convolutional processing.

# Training definitions

- Minimax
- Maximin
- Heuristic, non-saturating game
- Max likelihood game

# Minimax Game

- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$



- $D(x) = 1 \rightarrow$  The discriminator believes that  $x$  is a true image
- $D(G(z)) = 1 \rightarrow$  The discriminator believes that  $G(z)$  is a true image
- Equilibrium is a saddle point of the discriminator loss
- Final loss resembles Jensen-Shannon divergence

<https://arxiv.org/pdf/1701.00160.pdf>

# Minimax Game

- For the simple case of zero-sum game

$$J^{(G)} = -J^{(D)}$$

- So, we can summarize game by

$$V(\theta^{(D)}, \theta^{(G)}) = -J^{(D)}(\theta^{(D)}, \theta^{(G)})$$

- Easier theoretical analysis



# Minimax Game

- For the simple case of zero-sum game

$$J^{(G)} = -J^{(D)}$$

- So, we can summarize game by

$$V(\theta^{(D)}, \theta^{(G)}) = -J^{(D)}(\theta^{(D)}, \theta^{(G)})$$

- Easier theoretical analysis
- In practice not used → when the discriminator starts to recognize fake samples, then generator gradients vanish

# Heuristic non-saturating game

- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$
- $J^{(G)} = -\frac{1}{2} \mathbb{E}_{z \sim p_z} \log(D(G(z)))$
- Equilibrium not any more describable by single loss
- Generator maximizes the log-probability of the discriminator being mistaken
  - Good  $G(z) \rightarrow D(G(z)) = 1 \rightarrow J^{(G)}$  is maximized
- Heuristically motivated; generator can still learn even when discriminator successfully rejects all generator samples

# Original Algorithm

**for** number of training iterations **do**

**for**  $k$  steps **do**

    Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$

    Sample minibatch of  $m$  real samples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$

    Update the discriminator  $D$  by stochastic gradient ascend.

    Gradient:

$$\frac{\partial}{\partial D} \left( \frac{1}{m} \sum_{i=1}^m \log D(\mathbf{x}^{(i)}) + \frac{1}{m} \sum_{i=1}^m \log(1 - D(G(\mathbf{z}^{(i)}))) \right) .$$

**end for**

    Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$


    Update the generator  $G$  by stochastic gradient ascend.

    Gradient:


$$\frac{\partial}{\partial G} \left( \frac{1}{m} \sum_{i=1}^m \log(D(G(\mathbf{z}^{(i)}))) \right) .$$

**end for**

# DCGAN Architecture



# Examples *[up to 2015]*



Man  
with  
glasses




Man



Woman

- + =



Woman with  
glasses

# Modifying GANs for Max-Likelihood

- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$
- $J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log(\sigma^{-1}(D(G(z))))$

When discriminator is optimal, the generator gradient matches that of maximum likelihood


<https://arxiv.org/abs/1412.6515>



# Comparison of Generator Losses

When sample is likely fake, the minimax and the ML cost are flat  
→ no gradients in early steps

The ML cost variant generates gradients mostly from the “good generations”  
→ all gradients from few samples  
→ high variance  
→ Variance reduction?



# GAN Problems: Vanishing Gradients

- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_z \log(1 - D(G(z)))$
- $J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log(D(G(z)))$
- If the discriminator is quite bad
  - No accurate feedback for generator
  - No reasonable generator gradients
- But, if the discriminator is perfect,  $D(x) = D^*(x)$ 
  - Gradients go to 0
  - No learning anymore
- Bad when this happens early in the training
  - Easier to train the discriminator than the generator






Figure from Arjovsky and Bottou 2016. DCGAN after training 1, 10 and 25 epochs




# GAN Problems: Low dimensional supports

- Data lie in low-dim manifolds
- However, the manifold is not known
- During training  $p_g$  is not perfect either, especially in the start
- So, the support of  $p_r$  and  $p_g$  is non-overlapping and disjoint
- Easy to find a discriminating line



# GAN Problems: Batch Normalization does not work right way

- Batch-normalization causes strong intra-batch correlation
  - Activations depend on other inputs
  - Generations depend on other inputs
- Generation looks smooth but awkward, strong intra batch correlation



# Reference Batch Normalization

- Fix a reference batch  $R = \{r_1, r_2, \dots, r_m\}$
- Given new inputs  $X = \{x_1, x_2, \dots, x_m\}$
- Compute mean and standard deviation of feature of  $R$
- Normalize the features of  $X$  using the mean and standard deviation from  $R$
- Every  $x_1$  is always treated the same, regardless of which other examples appear in the minibatch

# Visual Batch Normalization

- Reference batch norm can overfit to the reference batch. A partial solution is virtual batch norm
- Fix a reference batch  $R = \{r_1, r_2, \dots, r_m\}$
- Given new inputs  $X = \{x_1, x_2, \dots, x_m\}$
- For each  $x_i$  in  $X$ :
  - Construct a virtual batch  $V$  containing both  $x_i$  and all of  $R$
  - Compute mean and standard deviation of features of  $V$
  - Normalize the features of  $x_i$  using the mean and standard deviation from  $V$

# Balancing generator and discriminator



- Usually the discriminator wins
  - Good, as the theoretical justification assumes a perfect discriminator
- Usually the discriminator network is bigger and deeper than the generator
- Sometimes running discriminator more often than generator works better
  - However, no real consensus
- Do not limit the discriminator to avoid making it too smart
  - Making learning “easier” will not necessarily make generation better
  - Better use non-saturating cost
  - Better use label smoothing

# Challenge: Convergence

- Optimization is tricky and unstable
  - Finding a saddle point does not imply a global minimum
  - A saddle point is also sensitive to disturbances
- An equilibrium might not even be reached
- Mode-collapse is the most severe form of non-convergence

# GAN Problems: Mode collapse


- Discriminator converges to the correct distribution
- Generator however places all mass in the most likely point
- All other modes are ignored
  - Underestimating variance
- Low diversity in generating samples





# Minibatch features

- Classify each sample by comparing to other examples in the mini-batch
- If samples are too similar, the model is penalized



# Challenge: how to evaluate?

- Despite the nice images, who cares?
- It would be nice to quantitatively evaluate the model
- For GANs it is hard to even estimate the likelihood
- In the absence of a precise evaluation metric, do GANs do truly good generations or generations that appeal/ fool to the human eye?
  - Can we trust the generations for critical applications, like medical tasks?
  - *Are humans a good discriminator for the converges generator?*

# Training procedure

- Use SGD-like algorithm of choice
  - Adam Optimizer is a good choice
- Use two mini-batches simultaneously
  - The first mini-batch contains real examples from the training set
  - The second mini-batch contains fake generated examples from the generator
- Optional: run k-steps of one player (e.g. discriminator) for every step of the other player (e.g. generator)

# Feature matching

- Instead of matching image statistics, match feature statistics

$$J_D = \left\| \mathbb{E}_{x \sim p_{data}} f(x) - \mathbb{E}_{z \sim p(z)} f(G(z)) \right\|_2^2$$

- $f$  can be any statistic of the data, like the mean or the median

# Use labels if possible

- Learning a conditional model  $p(y|x)$  is often generates better samples
  - Denton et al., 2015
- Even learning  $p(x,y)$  makes samples look more realistic
  - Salimans et al., 2016
- Conditional GANs are a great addition for learning with labels

# Summary

- GANs are generative models using supervised learning to approximate and intractable cost function
- GANs can simulate many cost functions, including max likelihood
- Finding Nash equilibria in high-dimensional, continuous, non-convex games is an important open research problem
- GAN research is in its infancy, most works published only in 2016. Not mature enough yet, but very compelling results



2014



2015



2016



2017



2018


# Models

- Wasserstein GAN
- Progressive GANs
- InfoGAN
- Conditional GAN
- StyleGAN
- CycleGAN



# Wasserstein GAN [Intuition]

- The distribution of the generated data should be as close as possible to the distribution of the real data [M. Arjovsky et al. 2017]
- The Wasserstein metric is the cost of optimal transport between the two distributions.




Intuitive (but imperfect) view:

$$W(\mathbb{P}_r, \mathbb{P}_\theta) = \min_{\gamma \in \Gamma} \sum_{(x,y) \in \gamma} [\|x - y\|],$$

where  $\Gamma$  is the set of all possible sets of correspondences between  $x$  and  $y$ .

# Progressive GANs

- T. Karras et al. 2018




# Progressive GANs: Results

Generated  
Image

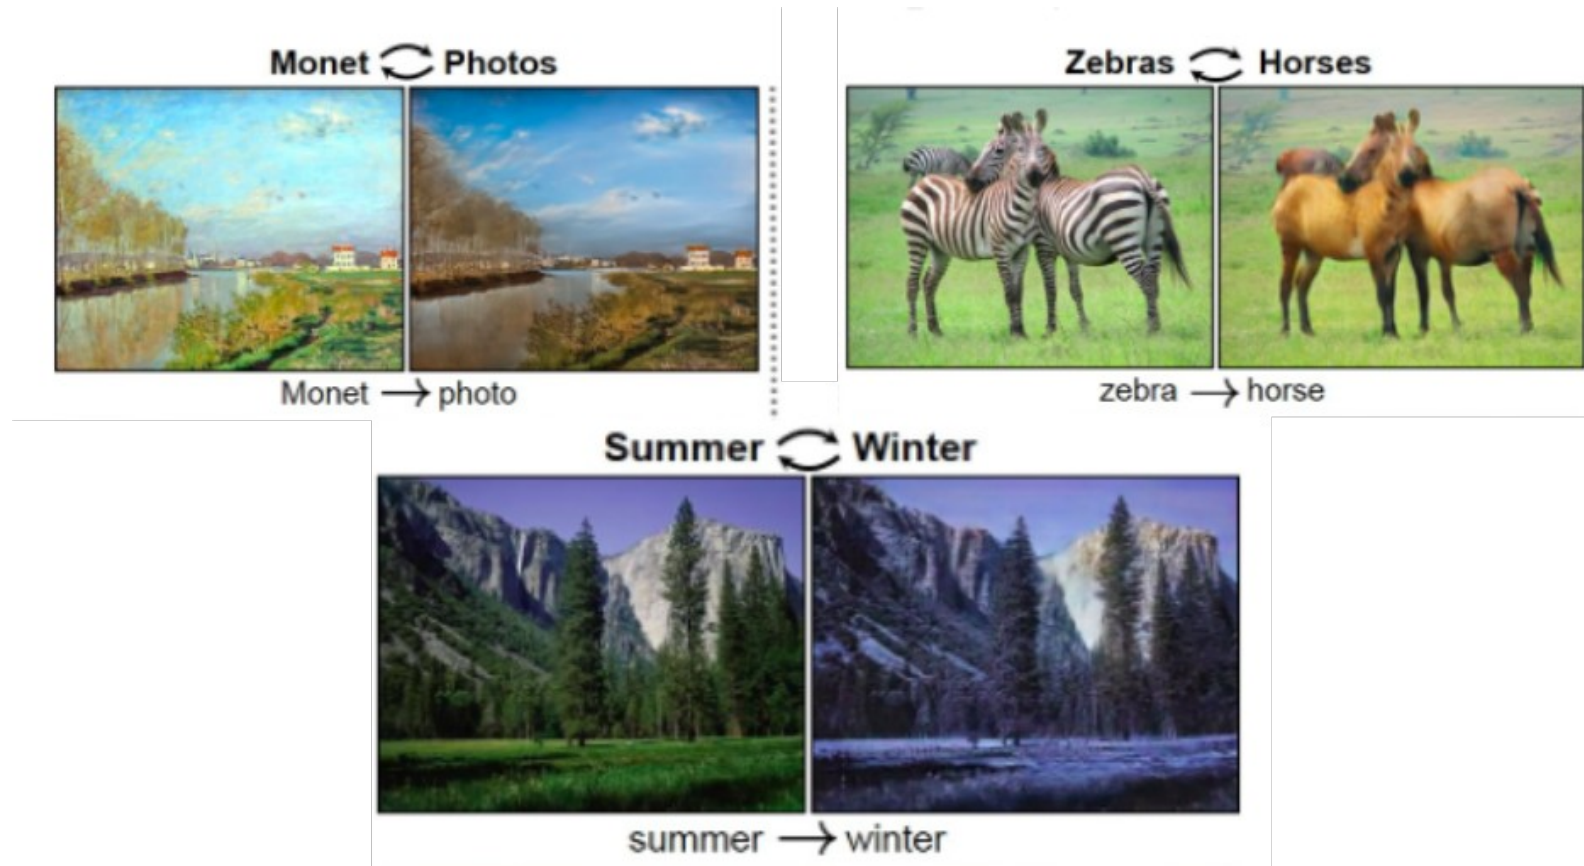


Nearest  
Neighbor  
in the  
training set




# CycleGAN

- How can we make sure we preserve the content of an input image?



# CycleGAN

- J.-Y. Zhu et al. 2017





# CycleGAN

- J.-Y. Zhu et al. 2017

